

Qcow2 research status

- The complexity of the performance bottleneck previously identified is actually linear – not factorial
→ does not change the problem
- tests on HDD vs SSD does not show change, the bottleneck is still there
- RAM consumption: it's the L2 indexation cache, 1 cache (2.5 MB) per snapshot in the chain

Qcow2 research status

- I have some ideas of ad-hoc solutions
 - perf: see the entire chain as a single disk and avoid iterating
 - memory (more interesting):